AI Ethics for Developers

Building Responsible AI Systems in Practice

2025-08-05

Webinar Overview

As AI systems become increasingly integrated into critical domains such as healthcare, finance, hiring, and criminal justice, the ethical implications of these technologies cannot be overlooked. Developers play a crucial role in ensuring AI systems are built responsibly from the ground up.

In this comprehensive webinar, Dr. Sarah Wilkins, a renowned AI ethics researcher, will bridge the gap between theoretical ethical principles and practical implementation considerations for developers. Participants will gain concrete strategies for identifying potential ethical issues, mitigating bias, ensuring fairness, and building transparent and accountable AI systems.

Key Topics

- **Ethical Frameworks for AI Development**: Understanding key principles and their practical applications
- **Bias Identification and Mitigation**: Techniques for detecting and addressing bias in data and models
- **Fairness Metrics and Trade-offs**: Quantifiable approaches to measuring and balancing different notions of fairness
- **Privacy-Preserving Techniques**: Methods for building AI systems that respect user privacy
- **Explainability and Interpretability**: Tools and approaches for making AI systems more transparent
- **Testing for Ethical Concerns**: Developing test suites to identify potential ethical issues
- **Documentation Practices**: Creating comprehensive documentation for datasets and models
- **Responsible Deployment**: Strategies for monitoring systems and addressing issues in production

Agenda

Time	Topic
00:00 - 15:00	Introduction to AI ethics and key principles
15:00 - 35:00	Data ethics: collection, representation, and bias
35:00 - 55:00	Algorithmic fairness: metrics, trade-offs, and implementation
55:00 - 70:00	Privacy considerations in AI development
70:00 - 90:00	Transparency, explainability, and documentation
90:00 - 105:00	Practical implementation strategies and case studies
105:00 - 120:00	Q&A and discussion

About the Presenter

Dr. Sarah Wilkins is a leading researcher in AI ethics with a focus on developing practical tools and methodologies for ethical AI implementation. She currently leads the AI Ethics Research division at the Institute for Responsible Technology, where she works with industry partners to develop ethical guidelines and assessment frameworks.

Dr. Wilkins holds a Ph.D. in Computer Science with a specialization in AI ethics from Stanford University and has published numerous papers on fairness, accountability, and transparency in AI systems. Her recent book, "Ethical AI by Design," has become a standard reference for practitioners seeking to implement responsible AI practices.

Who Should Attend

- Software engineers and developers working on AI systems
- Data scientists building and deploying machine learning models
- AI/ML researchers concerned with ethical implications of their work
- Engineering managers overseeing AI development teams
- Product managers guiding AI product development
- IT leaders responsible for AI strategy and governance
- Ethics officers and compliance professionals in tech organizations

Registration Information

Registration for this webinar opens on June 15, 2025. Due to the interactive nature of this session, attendance will be limited to 200 participants.

- **Standard Registration**: $35 (includes access to the live session and recording)
- **Premium Registration**: $75 (includes additional reference materials, code examples, and a follow-up Q&A session)

Frequently Asked Questions

What programming languages or frameworks will be covered?

The webinar will include examples in Python using common ML frameworks like TensorFlow and PyTorch, but the principles and approaches discussed are applicable across languages and frameworks.

Is this webinar suitable for beginners in AI/ML?

A basic understanding of machine learning concepts is recommended, but deep technical expertise is not required. The focus is on practical approaches that can be applied at various levels of technical proficiency.

Will there be coding demonstrations?

Yes, the webinar includes practical code examples demonstrating implementation of fairness metrics, bias testing, and explainability tools.

How does this differ from philosophical discussions about AI ethics?

While we'll cover key ethical frameworks, this webinar emphasizes practical implementation rather than theoretical discussions. You'll leave with concrete techniques and tools you can apply immediately in your work.